

1.FUN WITH VECTOR CALCULUS

1. Assume we have n real-valued scalar data points x_1, x_2, \dots, x_n . Analytically derive constant μ for which, $\sum_{i=1}^n (x_i - \mu)^2$ is minimized.

Here we set it as the loss function $L(\mu)$

$$L(\mu) = \sum_{i=1}^n (x_i - \mu)^2$$

In order to get the minimum, we set $\nabla L(\mu) = 0$

$$\nabla L(\mu) = 2 \sum_{i=1}^n (x_i - \mu) = 0$$

$$\sum_{i=1}^n (x_i - \mu) = 0$$

$$\sum_{i=1}^n x_i - n\mu = 0$$

$$\mu = \frac{1}{n} \sum_{i=1}^n x_i$$

Thus, when μ is the mean of x_1, x_2, \dots, x_n , the loss function is minimized.

2. Assume we have n data points are real d-dimensional vectors. Analytically derive a constant μ for which, $\sum_{i=1}^n \|x_i - \mu\|_2^2$ is minimized.

We can easily know that μ is also a d-dimensional vector. Here we set it as the loss function $L(\mu)$

$$L(\mu) = \sum_{i=1}^n \|x_i - \mu\|_2^2 = \sum_{i=1}^n \sum_{j=1}^d \sqrt{(x_{ij} - \mu_j)^2}^2 = \sum_{i=1}^n \sum_{j=1}^d (x_{ij} - \mu_j)^2$$

In order to get the minimum, we need to let the derivative of every $\mu_j (j \in (1, \dots, d))$ equals to zero so as to get the derivative of μ equals to zero.

$$\frac{\partial L(\mu)}{\partial \mu_j} = -2 \sum_{i=1}^n (x_{ij} - \mu_j) = 0$$

$$\mu_j = \frac{1}{n} \sum_{i=1}^n x_{ij}$$

Thus, we can get the value of μ ,

$$\mu = \begin{bmatrix} \frac{\sum_{i=1}^n x_{i1}}{n} \\ \frac{\sum_{i=1}^n x_{i2}}{n} \\ \vdots \\ \frac{\sum_{i=1}^n x_{id}}{n} \end{bmatrix}$$

Thus, when μ is the mean vector of x_1, x_2, \dots, x_n , the function is minimized.

2.LINEAR REGRESSION WITH NON-STANDARD LOSSES

1. Using matrix/vector notation, write down a loss function that measures the training error in terms of the l_1 -norm.

Assuming there are m data points. For each data point x , the number of its features is n :

$$x^{(i)} = \begin{bmatrix} x_0^{(i)} \\ x_1^{(i)} \\ \vdots \\ x_n^{(i)} \end{bmatrix}$$

For every data point, the $x_0 = 1$ so as to serve the bias. From the question we know that X is the matrix of training data points (stacked row-wise)

$$X = \begin{bmatrix} x_0^{(1)T} \\ x_1^{(1)T} \\ \vdots \\ x_m^{(1)T} \end{bmatrix} = \begin{bmatrix} 1 & x_1^{(1)} & x_2^{(1)} & \cdots & x_n^{(1)} \\ 1 & x_1^{(2)} & x_2^{(2)} & \cdots & x_n^{(2)} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 1 & x_1^{(m)} & x_2^{(m)} & \cdots & x_n^{(m)} \end{bmatrix}$$

For each data point x_i , it has a label y_i

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix}$$

We set the weight as ω :

$$\omega = \begin{bmatrix} \omega_0 \\ \omega_1 \\ \vdots \\ \omega_m \end{bmatrix}$$

So the predictions of the data points should be:

$$\begin{bmatrix} \omega_0 x_0^{(1)} + \omega_1 x_1^{(1)} + \cdots + \omega_n x_n^{(1)} \\ \vdots \\ \omega_0 x_0^{(m)} + \omega_1 x_1^{(m)} + \cdots + \omega_n x_n^{(m)} \end{bmatrix} = X\omega$$

Thus the loss should be the l_1 -norm difference between labels Y and predictions:

$$L(\omega) = \sum_{i=1}^m |y_i - \sum_{j=1}^n \omega_j x_j^{(i)}| = \sum_{i=1}^m |y_i - \omega^T x^{(i)}|$$

For conciseness, we write this as:

$$L(\omega) = \|X\omega - y\|_1$$

2. Can you write down the optimal linear model in closed form? If not, why not? So we can't write down the optimal linear model in closed form.

No, because the l_1 loss function is not a continuous derivative function.

From the Figure 1, we can easily know that when $X\omega - Y = 0$, the loss is the minima. However, this point doesn't have derivative. Thus, we cannot find the minima by finding the point whose derivative is zero.

3. If the answer to b is no, can you think of an alternative algorithm to optimize the loss function? Comment on its pros and cons.

Although using gradient descent we can get a comparatively small value closed to the optimal of l_1 loss function, it might be hard to reach the optimal(minima). The gradient of the l_1 loss



Figure 1. L1 loss function

function only have two values, thus when it is in neighborhood of the optimal, it might still have a comparatively large gradient, thus it might miss the optimal and jump around the minima.

However, in other words, if we try to minimize l_1 loss (a.k.a MAE), that prediction would be the median of all observations. Here's how to prove it:

Assuming we have y_1, y_2, \dots, y_n and β is the prediction.

$$L_1 = \sum_{i=1}^n |y_i - \beta|$$

$$\frac{\partial L_1}{\partial \beta} = - \sum_{i=1}^n \text{sgn}(y_i - \beta)$$

$$\text{sgn}(y_i - \beta) = \begin{cases} 1 & \text{if } y_i > \beta \\ -1 & \text{if } y_i < \beta \end{cases}$$

Thus, the derivative equals to 0 when there is the same number of positive and negative terms among the $y_i - \beta$, which means β should be the median of y_i

Pros:

- a. More robust to outliers compared to L2 loss;
- b. Easy to understand and implement.

Cons:

- a. The complexity for finding median would be high when the data set is huge;
- b. Median can be a bit jumpy in small samples made up of discrete values;
- c. When data points are a bit of clustered, using median might not a good choice.

3.HARD CODING A MULTI-LAYER PERCEPTRON

The multi-layer perceptron network structure is shown in figure 2:

for Perceptron1 P_1 :

$$f(x) = \begin{cases} 1 & \text{if } x_2 - x_1 > 0 \\ 0 & \text{otherwise} \end{cases}$$

for Perceptron1 P_2 :

$$f(x) = \begin{cases} 1 & \text{if } x_3 - x_2 > 0 \\ 0 & \text{otherwise} \end{cases}$$




Figure 2. structure of the neural network.

for Perceptron1 P_3 :

$$f(x) = \begin{cases} 1 & \text{if } x_4 - x_3 > 0 \\ 0 & \text{otherwise} \end{cases}$$

for Perceptron1 P_4 :

$$f(x) = \begin{cases} 1 & \text{if } x_5 - x_4 > 0 \\ 0 & \text{otherwise} \end{cases}$$

So only when all perceptrons have the output as 1, then it meets the requirement that $x_1 < x_2 < x_3 < x_4 < x_5$. Thus, for Perceptron out P_{out} :

$$f(x) = \begin{cases} 1 & \text{if } x_1 + x_2 + x_3 + x_4 = 4 \\ -1 & \text{otherwise} \end{cases}$$

OK, thus far we have been talking about linear models. All these can be viewed as a single-layer neural net. The next step is to move on to multi-layer nets. Training these is a bit more involved, and implementing from scratch requires time and effort. Instead, we just use well-established libraries. I prefer PyTorch, which is based on an earlier library called Torch (designed for training neural nets via backprop).

```
In [ ]: import numpy as np
import torch
import torchvision
```

Torch handles data types a bit differently. Everything in torch is a *tensor*.

```
In [ ]: a = np.random.rand(2,3)
b = torch.from_numpy(a)

# Q4.1 Display the contents of a, b
print("contest of a: ", a)
print("contest of b: ", b)

contest of a: [[0.816508  0.75597224 0.67269371]
 [0.38200301 0.06110529 0.70656539]]
contest of b: tensor([[0.8165, 0.7560, 0.6727],
 [0.3820, 0.0611, 0.7066]]), dtype=torch.float64)
```

The idea in Torch is that tensors allow for easy forward (function evaluations) and backward (gradient) passes.

```
In [ ]: A = torch.rand(2,2)
b = torch.rand(2,1)
x = torch.rand(2,1, requires_grad=True)

y = torch.matmul(A,x) + b

print(y)
z = y.sum()
print(z)
z.backward()
print(x.grad)
print(x)
```

```
tensor([[0.9157],
        [1.3327]]), grad_fn=<AddBackward0>)
tensor(2.2484, grad_fn=<SumBackward0>)
tensor([[0.7845],
        [1.7336]])
tensor([[0.1465],
        [0.9142]]), requires_grad=True)
```

Notice how the backward pass computed the gradients using autograd. OK, enough background. Time to train some networks. Let us load the *Fashion MNIST* dataset, which is a database of grayscale images of clothing items.

```

In [ ]: trainingdata = torchvision.datasets.FashionMNIST('./FashionMNIST/', train=True, download=True, transform=
=torchvision.transforms.ToTensor())
testdata = torchvision.datasets.FashionMNIST('./FashionMNIST/', train=False, download=True, transform=to
rchvision.transforms.ToTensor())

Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-images-idx3-ubyte.gz to
./FashionMNIST/FashionMNIST/raw/train-images-idx3-ubyte.gz

Extracting ./FashionMNIST/FashionMNIST/raw/train-images-idx3-ubyte.gz to ./FashionMNIST/FashionMNIS
T/raw
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/train-labels-idx1-ubyte.gz to
./FashionMNIST/FashionMNIST/raw/train-labels-idx1-ubyte.gz

Extracting ./FashionMNIST/FashionMNIST/raw/train-labels-idx1-ubyte.gz to ./FashionMNIST/FashionMNIS
T/raw
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-images-idx3-ubyte.gz to
./FashionMNIST/FashionMNIST/raw/t10k-images-idx3-ubyte.gz

Extracting ./FashionMNIST/FashionMNIST/raw/t10k-images-idx3-ubyte.gz to ./FashionMNIST/FashionMNIST/
raw
Downloading http://fashion-mnist.s3-website.eu-central-1.amazonaws.com/t10k-labels-idx1-ubyte.gz to
./FashionMNIST/FashionMNIST/raw/t10k-labels-idx1-ubyte.gz

Extracting ./FashionMNIST/FashionMNIST/raw/t10k-labels-idx1-ubyte.gz to ./FashionMNIST/FashionMNIST/
raw
Processing...
Done!

/usr/local/lib/python3.6/dist-packages/torchvision/datasets/mnist.py:469: UserWarning: The given Num
Py array is not writeable, and PyTorch does not support non-writeable tensors. This means you can wr
ite to the underlying (supposedly non-writeable) Numpy array using the tensor. You may want to copy
the array to protect its data or make it writeable before converting it to a tensor. This type of wa
rning will be suppressed for the rest of this program. (Triggered internally at /pytorch/torch/csr
c/utils/tensor_numpy.cpp:141.)
return torch.from_numpy(parsed.astype(m[2], copy=False)).view(*s)

```

Let us examine the size of the dataset.

```
In [ ]:
```

```
# Q4.2 How many training and testing data points are there in the dataset?  
# What is the number of features in each data point?
```


```
# There are 60000 training data points and 10000 testing data points in the dataset  
# For each data point, we have the image data and the actual label for this image.  
# The data size of each image is 1*28*28  
# '1' means each pixel only has one channel, that's grayscale value  
# '28*28' means the width and height of each image are both 28 pixels,  
# that means each image has 28*28=784 pixels  
print("the number of training data points: ", len(trainingdata))  
print("the number of testing data points: ", len(testdata))  
print("the size of each data point vector: ", len(trainingdata[0]))  
print("the size of each image data: ", trainingdata[0][0].size())  
print("label for each data is", type(trainingdata[0][1]))
```

```
the number of training data points: 60000  
the number of testing data points: 10000  
the size of each data point vector: 2  
the size of each image data: torch.Size([1, 28, 28])  
label for each data is <class 'int'>
```

Let us try to visualize some of the images. Since each data point is a tensor (not an array) we need to postprocess to use matplotlib.


```
In [ ]: import matplotlib.pyplot as plt
%matplotlib inline

image, label = trainingdata[0]
# Q4.3 Assuming each sample is an image of size 28x28, show it in matplotlib.
plt.figure()
plt.imshow(image[0], cmap='gray')
plt.show()
```



Let's try plotting several images. This is conveniently achieved in PyTorch using a *data loader*, which loads data in batches.


```
In [ ]: trainDataLoader = torch.utils.data.DataLoader(trainingdata, batch_size=64, shuffle=True)
testDataLoader = torch.utils.data.DataLoader(testdata, batch_size=64, shuffle=False)
images, labels = iter(trainDataLoader).next()
print(images.size(), labels)

torch.Size([64, 1, 28, 28]) tensor([5, 4, 6, 6, 6, 2, 1, 7, 1, 8, 4, 8, 8, 1, 6, 8, 8, 0, 8, 7, 4, 2,
      9, 0, 7,
      2, 2, 1, 3, 8, 2, 2, 4, 0, 1, 4, 5, 3, 3, 2, 5, 5, 7, 3, 6, 4, 1, 5, 2,
      4, 7, 1, 0, 7, 1, 4, 6, 5, 0, 2, 6, 2, 7, 2, 6])
```

```
In [ ]: # Q4.4 Visualize the first 10 images of the first minibatch
# returned by testDataloader.
```

```
first_ten_images = images[:10, 0, ...].numpy()
row = np.concatenate([first_ten_images[i] for i in range(10)], axis=1)
plt.figure()
print("Showing the 10 images in one row:")
plt.imshow(row, cmap='gray')
plt.show()
```

Showing the 10 images in one row:



Now we are ready to define our linear model. Here is some boilerplate PyTorch code that implements the forward model for a single layer network for logistic regression (similar to the one discussed in class notes).

```
In [ ]: class LinearReg(torch.nn.Module):
def __init__(self):
super(LinearReg, self).__init__()
self.linear = torch.nn.Linear(28*28, 10)

def forward(self, x):
x = x.view(-1, 28*28)
transformed_x = self.linear(x)
return transformed_x

net = LinearReg().cuda()
loss = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(net.parameters(), lr=0.01)
```

Cool! Now we have set everything up. Let's try to train the network.

```
In [ ]: train_loss_history = []
        test_loss_history = []

# Q4.5 Write down a for-loop that trains this network for 20 minibatch iterations,
# and print the train/test losses.
# Save them in the variables above. If done correctly, you should be able to
# execute the next code block.

#Initiate network
net = LinearReg().cuda()
loss = torch.nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(net.parameters(), lr=0.01)

num_of_epochs = 20

for i in range(num_of_epochs):
    #shuffle training data
    trainDataLoader = torch.utils.data.DataLoader(trainingdata, batch_size=64, shuffle=True)

    #train the net with each minibatch
    epoch_train_loss = 0
    for images, labels in trainDataLoader:
        optimizer.zero_grad()
        # calculate train loss of the current minibatch
        preds = net(images.cuda())
        batch_images_loss = Loss(preds, labels.cuda())
        epoch_train_loss += batch_images_loss

    #optimize the network
    batch_images_loss.backward()
    optimizer.step()

    # calculate train loss for the epoch
    epoch_train_loss /= len(trainingdata)
    train_loss_history.append(epoch_train_loss)

    # test the net with each minibatch
    epoch_test_loss = 0
    for images, labels in testDataLoader:
        preds = net(images.cuda())
        batch_images_loss = Loss(preds, labels.cuda())
        epoch_test_loss += batch_images_loss
```


```
# calculate test loss for the epoch
epoch_test_loss /= len(testdata)
test_loss_history.append(epoch_test_loss)
```

```
print('Epoch: {}, Train Loss: {}, Test Loss: {}'.format(i, epoch_train_loss, epoch_test_loss))
```

```
Epoch: 0, Train Loss: 0.01502606924623251, Test Loss: 0.011578425765037537
Epoch: 1, Train Loss: 0.010432981885969639, Test Loss: 0.010125137865543365
Epoch: 2, Train Loss: 0.009428427554666996, Test Loss: 0.009466097690165043
Epoch: 3, Train Loss: 0.008885594084858894, Test Loss: 0.00906306505203247
Epoch: 4, Train Loss: 0.008530200459063053, Test Loss: 0.008839458227157593
Epoch: 5, Train Loss: 0.008271351456642151, Test Loss: 0.008571295067667961
Epoch: 6, Train Loss: 0.00807273667305708, Test Loss: 0.00842297449707985
Epoch: 7, Train Loss: 0.007911734282970428, Test Loss: 0.008297959342598915
Epoch: 8, Train Loss: 0.007779350504279137, Test Loss: 0.008169146254658699
Epoch: 9, Train Loss: 0.007666518911719322, Test Loss: 0.00807636696845293
Epoch: 10, Train Loss: 0.007574894465506077, Test Loss: 0.008048789575695992
Epoch: 11, Train Loss: 0.007492510136216879, Test Loss: 0.007952781394124031
Epoch: 12, Train Loss: 0.007416723761707544, Test Loss: 0.007868158631026745
Epoch: 13, Train Loss: 0.0073526217602193356, Test Loss: 0.0078074736520648
Epoch: 14, Train Loss: 0.007288788910955191, Test Loss: 0.007769959978759289
Epoch: 15, Train Loss: 0.007235296536237001, Test Loss: 0.007714950945228338
Epoch: 16, Train Loss: 0.0071831876412034035, Test Loss: 0.0077091665007174015
Epoch: 17, Train Loss: 0.007145352195948362, Test Loss: 0.007648747880011797
Epoch: 18, Train Loss: 0.007103899493813515, Test Loss: 0.007626079488545656
Epoch: 19, Train Loss: 0.007061449345201254, Test Loss: 0.007573876064270735
```

```
In [ ] : plt.plot(range(20),train_loss_history,'-',linewidth=3,label='Train error')
plt.plot(range(20),test_loss_history,'-',linewidth=3,label='Test error')
plt.xlabel('epoch')
plt.ylabel('loss')
plt.grid(True)
plt.legend()
```

Out [] : <matplotlib.legend.Legend at 0x7ff243392e710>



Neat! Now let's evaluate our model accuracy on the entire dataset. The predicted class label for a given input image can be computed by looking at the output of the neural network model and computing the index corresponding to the maximum activation. Something like

```
predicted_output = net(images) → predicted_labels = torch.max(predicted_output,1)
```

```
In [ ]: predicted_output = net(images.cuda())
print(torch.max(predicted_output, 1))
fit = Loss(predicted_output, labels.cuda())
print(labels)
```

```
torch.return_types.max(
values=tensor([ 6.1149,  2.9275,  8.9590,  7.8504,  6.6379,  6.2165, 10.2046,  4.3740,
 7.0905, 11.6778, 10.4770, 10.4350,  6.6251,  4.7262,  9.3121,  4.5606],
device='cuda:0', grad_fn=<MaxBackward0>),
indices=tensor([3, 1, 7, 5, 8, 2, 5, 6, 8, 9, 1, 9, 1, 8, 1, 5], device='cuda:0'))
tensor([3, 2, 7, 5, 8, 4, 5, 6, 8, 9, 1, 9, 1, 8, 1, 5])
```

```
In [ ]: def evaluate(data_loader):
```

```
    # Q4.6 Implement a function here that evaluates training and testing accuracy.
    # Here, accuracy is measured by probability of successful classification.
```

```
    correct_pred = 0
    num_of_all_images = 0
    for images, labels in data_loader:
        predicted_output = net(images.cuda())
        _, predicted_labels = torch.max(predicted_output, 1)
        preds_diff = labels - predicted_labels.cpu()
        correct_pred += int((preds_diff==0).sum())
        num_of_all_images += len(images)
    return correct_pred/num_of_all_images
```

```
print('Train acc = %0.2f, test acc = %0.2f' % (evaluate(trainDataLoader), evaluate(testDataLoader)))
```

```
Train acc = 0.85, test acc = 0.83
```